TELOS: Mathematical Enforcement of AI Constitutional Boundaries

Achieving Near-Zero Attack Success Rate Through Embedding-Space Governance

Jeffrey Brunner
TELOS AI Labs Inc.
JB@telos-labs.ai

January 2026

Abstract

We present TELOS, a runtime AI governance system that achieves a 0% observed Attack Success Rate (ASR) across 1,300 adversarial attacks (95% CI: [0%, 0.28%]). While current systems accept violation rates of 3.7% to 43.9% as unavoidable, TELOS demonstrates that mathematical enforcement of constitutional boundaries can provide substantially stronger defense than existing approaches. It uses a three-tier structure that combines embedding-space mathematics, authoritative policy retrieval, and human expert escalation.

Our approach treats constitutional enforcement as a statistical process control problem rather than a prompt engineering challenge. We use fixed reference points in embedding space (Primacy Attractors) combined with control-theoretic stability analysis to create a governance framework that achieved strong results against tested attack vectors.

We validate our method across 1,300 attacks (400 from HarmBench general-purpose and 900 from MedSafetyBench healthcare-specific), covering various harm categories from direct violations to complex jailbreaks. TELOS-governed models achieve 0% ASR on both small and large language models. In contrast, baseline methods using system prompts show an ASR of 3.7–11.1%, while raw models exhibit an ASR of 30.8–43.9%.

The system includes governance trace logging that makes enforcement decisions observable and auditable, supporting regulatory compliance requirements. All results are fully reproducible with the provided code and attack libraries.

Keywords: AI safety, constitutional AI, adversarial robustness, embedding space, Lyapunov stability, governance verification

1 Introduction

The deployment of Large Language Models (LLMs) in regulated fields such as healthcare, finance, and education presents a fundamental conflict. These systems offer significant capabilities, but they lack reliable ways to enforce regulatory boundaries. This conflict has become legally urgent. The European Union's AI Act requires runtime monitoring and ongoing compliance for high-risk AI systems by August 2026. Meanwhile, California's SB 243 is the first state law aimed explicitly at AI chatbot safety for minors, effective January 2026. These regulations demand mechanisms that current governance approaches cannot provide.

Current methods for AI governance, whether through fine-tuning, prompt engineering, or post-hoc filtering, often fail against adversarial attacks. The HarmBench benchmark found that leading AI systems show attack success rates of 4.4–90% across 400 standardized attacks. MedSafetyBench revealed similar weaknesses in healthcare contexts with 900 domain-specific attacks. Leading guardrail systems, such as NVIDIA NeMo Guardrails and Llama Guard, accept violation rates between 3.7% and 43.9% as unavoidable, which is incompatible with emerging regulatory requirements.

This paper investigates whether substantially lower failure rates are achievable through a different architectural approach. We apply statistical process control methods to AI governance and present empirical evidence that constitutional enforcement can be significantly strengthened.

1.1 The Governance Problem

Consider a healthcare AI assistant that must never disclose Protected Health Information (PHI) under HIPAA regulations. Current methods fail in predictable ways:

1. **Prompt Engineering:** System prompts stating "never disclose PHI" can easily be bypassed using social engineering or prompt injection.
2. **Fine-tuning:** RLHF/DPO methods embed constraints into model weights but remain vulnerable to jailbreaks.
3. **Output Filtering:** Filtering after generation captures obvious violations but overlooks semantic equivalents.

The core issue is that all current methods treat governance as a linguistic problem (what the model states) rather than a geometric problem (the location of the query in semantic space).

1.2 Our Approach: Governance as Geometric Control

TELOS implements AI governance through three architectural choices:

1. **Fixed Reference Points:** Instead of relying on the model's shifting attention for self-governance, we set fixed reference points (Primacy Attractors) in the embedding space.
2. **Mathematical Enforcement:** Cosine similarity in the embedding space offers a deterministic, position-invariant measure of constitutional alignment.
3. **Three-Tier Defense:** The system ensures that mathematical (PA), authoritative (RAG), and human (Expert) layers must all fail simultaneously for a violation to occur.

1.3 Contributions

This paper makes four main contributions:

1. **Theoretical:** We demonstrate that external reference points in the embedding space enable stable governance with defined basin geometry ($r = 2/\rho$).
2. **Empirical:** We show 0% ASR across 1,300 adversarial attacks (400 HarmBench + 900 MedSafetyBench), compared to 3.7–43.9% for existing methods.
3. **Methodological:** We provide governance trace logging that enables forensic analysis and regulatory audit trails.
4. **Practical:** We provide reproducible validation scripts and a healthcare-specific implementation designed to support HIPAA compliance requirements (formal compliance requires independent audit and organizational safeguards beyond technical controls).

1.4 Threat Model

Our evaluation assumes a **query-only adversary** with the following characteristics:

- **Knowledge:** Attacker knows TELOS exists but not the specific PA configuration, threshold values, or embedding model details
- **Access:** Black-box query access only; no ability to modify embeddings, intercept API calls, or access system internals
- **Capabilities:** Can craft arbitrary text inputs, including multi-turn conversations, role-play scenarios, and prompt injection attempts
- **Limitations:** Cannot perform model extraction attacks, cannot modify the governance layer, and is subject to standard rate limiting

This threat model aligns with HarmBench and MedSafetyBench evaluation assumptions. We note that white-box adaptive attacks (where adversaries can optimize against known PA configurations) represent an important direction for future work, though such attacks require access levels uncommon in production deployments.

2 The Reference Point Problem

2.1 Why Attention Mechanisms Fail for Governance

Modern transformers use attention mechanisms to determine token relationships:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V \tag{1}$$

This creates a key problem for governance. The model generates both Q and K from its own hidden states, leading to self-referential circularity. Research on the "lost in the middle" effect [Liu et al., 2024] demonstrates that LLMs exhibit strong primacy and recency biases: they attend well to information at the beginning and end of context, but poorly to middle positions. As conversations extend, initial constitutional constraints drift into this poorly-attended middle region:

$$\text{Attention}(Q_i, K_j) \propto e^{-\alpha|i-j|} \quad \text{(simplified model)} \tag{2}$$

At position $i = 1000$, attention to initial constraints ($j = 0$) can drop substantially. The model effectively "forgets" its constitutional limits as context accumulates, a direct consequence of these documented positional attention biases.

2.2 The Primacy Attractor Solution

Instead of relying on self-reference, TELOS sets up an external, fixed reference point:

Definition (Primacy Attractor): A fixed point $\hat{a} \in \mathbb{R}^n$ in embedding space that includes constitutional constraints:

$$\hat{a} = \frac{\tau \cdot p + (1 - \tau) \cdot s}{\|\tau \cdot p + (1 - \tau) \cdot s\|} \tag{3}$$

Where:

- p = purpose vector (embedded purpose statements)
- s = scope vector (embedded boundaries)

- τ = constraint tolerance $\in [0, 1]$

The PA stays constant throughout conversations. This provides a stable reference for measuring fidelity:

$$\text{Fidelity}(q) = \cos(q, \hat{a}) = \frac{q \cdot \hat{a}}{\|q\| \cdot \|\hat{a}\|} \tag{4}$$

Note that because the PA encodes constitutional boundaries (prohibited behaviors), higher fidelity indicates a query closer to violation territory. This may seem counterintuitive, but it follows from the PA representing what must be blocked rather than what should be permitted.

This geometric relationship is independent of token position or context window, fixing the reference point problem.

3 Mathematical Foundation

3.1 Basin of Attraction

The basin $B(\hat{a})$ defines the area where queries align with the constitution:

Design Heuristic 1 (Basin Geometry): The basin radius is given by:

$$r = \frac{2}{\rho} \quad \text{where } \rho = \max(1 - \tau, 0.25) \tag{5}$$

Rationale: This formula is a geometric design heuristic chosen to balance false positives against adversarial coverage. The floor at $\rho = 0.25$ prevents unbounded basin growth; when tolerance is at maximum ($\tau = 0.9$), the basin radius is limited to $r = 8.0$, maintaining meaningful boundaries. We do not claim this formula is optimal; rather, it represents a principled design choice validated empirically through our benchmark testing.

3.2 Lyapunov Stability Analysis

We apply Lyapunov stability analysis from classical control theory to characterize the PA system. While embedding space operates in discrete inference steps rather than continuous time, the Lyapunov framework provides principled design intuition that we validate empirically in Section 5.

Definition (Lyapunov Function):

$$V(x) = \frac{1}{2}\|x - \hat{a}\|^2 \tag{6}$$

Proposition 2 (Global Asymptotic Stability): The PA system is globally stable with proportional control $u = -K(x - \hat{a})$ for $K > 0$.

Proof Sketch:

1. $V(x) = 0$ iff $x = \hat{a}$ (positive definite)
2. $\dot{V}(x) = \nabla V(x) \cdot \dot{x} = (x - \hat{a}) \cdot (-K(x - \hat{a})) = -K\|x - \hat{a}\|^2 < 0$ for $x \neq \hat{a}$
3. $V(x) \to \infty$ as $\|x\| \to \infty$ (radially unbounded)

By Lyapunov's theorem, these conditions establish global asymptotic stability for the idealized continuous dynamical system. The discrete, high-dimensional nature of embedding space means this analysis provides design guidance rather than formal guarantees; empirical validation (Section 5) confirms that the stability properties hold in practice.

3.3 Proportional Control Law

The intervention strength follows proportional control:

$$F(x) = K \cdot e(x) \quad \text{where } e(x) = \max(0, f(x) - \theta) \tag{7}$$

With $K = 1.5$ (empirically tuned) and threshold $\theta = 0.65$ (healthcare domain), this ensures:

- Immediate blocking for high-fidelity queries ($f \geq 0.65$) that closely match prohibited patterns
- Proportional correction for ambiguous drift ($0.35 \leq f < 0.65$)
- No Tier 1 intervention for low-fidelity queries ($f < 0.35$), though secondary heuristics may flag rare edge cases for Tier 3 review

4 Three-Tier Defense Architecture

TELOS uses defense-in-depth through three independent layers:

4.1 Tier 1: Mathematical Enforcement (Primacy Attractor)

- **Mechanism:** Embedding-based fidelity measurement
- **Decision:** Block if fidelity(query, PA) $\geq$ threshold
- **Properties:** Deterministic, position-invariant, millisecond latency

4.2 Tier 2: Authoritative Guidance (RAG Corpus)

- **Mechanism:** Retrieval-Augmented Generation from verified regulatory sources
- **Activation:** When $0.35 \leq$ fidelity < 0.65 (ambiguous zone)
- **Corpus:** Federal regulations (CFR), HIPAA guidance documents, professional standards (AMA, CDC)

Tier 2 addresses cases where mathematical similarity alone is insufficient: the query falls in an ambiguous zone. Rather than relying on the LLM's parametric knowledge (which may hallucinate policy), the system retrieves authoritative source text and grounds the response in documented regulations. This provides auditable, citation-backed reasoning for edge cases.

The RAG corpus is designed to grow over time. As the system encounters novel edge cases across deployments, domain-specific guidance can be incorporated, reducing the proportion of queries requiring Tier 2 or Tier 3 intervention. This continuous refinement follows the SPC principle of process improvement: each ambiguous case informs corpus expansion, progressively strengthening Tier 1 coverage.

4.3 Tier 3: Human Expert Escalation

- **Mechanism:** Domain experts with professional responsibility
- **Activation:** Edge cases where fidelity < 0.35 but secondary heuristics suggest potential novel attacks
- **Roles:** Privacy Officer, Legal Counsel, Chief Medical Officer

4.4 Low Probability of Simultaneous Failure

For a violation to occur, an attacker must simultaneously:

1. Manipulate embedding math (requires API access)
2. Exploit gaps in federal regulations (highly constrained)
3. Deceive trained professionals (unlikely under standard protocols)

The requirement that all three layers fail simultaneously makes successful attacks highly improbable under our evaluated threat model.

5 Adversarial Validation

5.1 Attack Taxonomy

We tested 1,300 attacks across two benchmarks:

Table 1: Attack Benchmark Summary

Benchmark	Source	Attacks	Domain	TELOS ASR
HarmBench	Center for AI Safety	400	General-purpose harms	0%
MedSafetyBench	NeurIPS 2024	900	Healthcare/medical safety	0%
Total		1,300		0%

Tier Distribution:
HarmBench ($n = 400$):

- Tier 1 (PA blocks): 95.8% (383/400)
- Tier 2 (RAG blocks): 3.0% (12/400)
- Tier 3 (Expert blocks): 1.2% (5/400)

MedSafetyBench ($n = 900$):

- Tier 1 (PA blocks): 23.0% (207/900)
- Tier 2 (RAG blocks): 77.0% (693/900)
- Tier 3 (Expert blocks): 0.0% (0/900)

The difference in tier distribution reflects the nature of the attacks: HarmBench attacks are more direct violations that trigger immediate Tier 1 blocking, while MedSafetyBench healthcare attacks often fall in the ambiguous zone requiring Tier 2 policy retrieval for resolution.

5.2 Experimental Setup

Models Tested:

- Mistral Small (baseline and TELOS-controlled)
- Mistral Large (baseline and TELOS-controlled)
- Raw models (no controls)

Metrics:

- Attack Success Rate (ASR): Percentage of successful violations
- Violation Defense Rate (VDR): $1 - \text{ASR}$

5.3 Results

Table 2: Attack Success Rate by Configuration

Configuration	ASR	VDR	95% CI
Raw Mistral Small	30.8%	69.2%	[25.1%, 36.5%]
Mistral Small + System Prompt	11.1%	88.9%	[7.8%, 14.4%]
Mistral Small + TELOS	0.0%	100.0%	[0.0%, 5.4%]
Raw Mistral Large	43.9%	56.1%	[37.8%, 50.0%]
Mistral Large + System Prompt	3.7%	96.3%	[1.9%, 5.5%]
Mistral Large + TELOS	0.0%	100.0%	[0.0%, 5.4%]

5.4 Statistical Significance

Using Wilson score intervals for 0 out of 1,300 successes:

- 95% CI: [0.0%, 0.28%]
- 99% CI: [0.0%, 0.35%]

This confirms a 0% ASR with high confidence. This result is significantly different from baseline approaches ($p < 0.001$, Fisher's exact test).

6 Limitations and Future Work

6.1 Current Limitations

This work represents an initial validation study with several limitations that require future investigation:

Model Coverage: All results use Mistral embeddings (Small and Large variants). We have not validated performance on other embedding models (OpenAI, Cohere, open-source alternatives) or other LLM families (GPT-4, Claude, Llama). Generalization to other models is an open question requiring additional validation.

Threat Model Scope: Our validation assumes black-box query access. We have not tested adaptive attacks where adversaries have knowledge of the PA configuration or can optimize against the embedding-space geometry. White-box robustness remains untested.

Domain Coverage: Validation focused on healthcare (MedSafetyBench) and general safety (HarmBench). Performance in other regulated domains (finance, legal, education) requires separate validation studies with domain-specific attack sets.

Language Coverage: All validation is English-only. Cross-lingual attacks and multilingual deployment scenarios are untested.

False Positive Analysis: We have not systematically measured false positive rates on benign conversations. Over-blocking in legitimate use cases could limit practical deployment.

Computational Overhead: Less than 50ms latency per query is acceptable for most applications but may not meet requirements for latency-critical deployments.

Human Scalability: Tier 3 expert escalation (1.2% of queries in our validation) does not scale to millions of daily queries without significant staffing infrastructure.

Multimodal: This work addresses text-only inputs. Image-based jailbreaks and multimodal attacks are out of scope.

6.2 Future Directions

1. **Multi-Modal Extension:** Expand PA to include image and audio inputs using CLIP-style embeddings.
2. **Adaptive PAs:** Implement federated learning for PA updates across consortium sites.
3. **Formal Verification:** Prove stronger properties beyond Lyapunov stability.
4. **Economic Analysis:** Conduct a cost-benefit study of TELOS versus manual compliance.

7 Conclusion

TELOS shows that AI constitutional violations are not unavoidable. Through three-tier governance (mathematical enforcement, authoritative policy retrieval, and human expert escalation) we achieve a 0% observed Attack Success Rate across 1,300 adversarial tests (95% CI: [0%, 0.28%]), with every threat detected and appropriately handled by the framework.

Our three contributions (theoretical: Lyapunov-stable PA mathematics; empirical: 0% ASR validation; methodological: governance trace logging for auditability) address a gap in current AI deployment practices for regulated fields.

The reference implementation is available for organizations seeking to evaluate the framework in their specific deployment contexts.

We invite the research community to reproduce and extend our findings. The code is open source (Apache 2.0) and the validation protocol is automated to support independent verification.

Extending this approach to additional models, domains, and threat models requires institutional collaboration and broader validation studies.

Data Availability

All validation datasets are available on Zenodo:

- Adversarial Validation: https://doi.org/10.5281/zenodo.18013104
- Governance Benchmark: https://doi.org/10.5281/zenodo.18009153
- SB 243 Child Safety: https://doi.org/10.5281/zenodo.18027446

Code repository: https://github.com/TelosSteward/TELOS

References

Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., and Liang, P. (2024). Lost in the Middle: How Language Models Use Long Contexts. *Transactions of the Association for Computational Linguistics.*

Mazeika, M., Phan, L., Yin, X., et al. (2024). HarmBench: A Standardized Evaluation Framework for Automated Red Teaming and Robust Refusal. *arXiv preprint arXiv:2402.04249.*

Han, T., Kumar, A., Agarwal, C., and Lakkaraju, H. (2024). MedSafetyBench: Evaluating and Improving the Medical Safety of Large Language Models. *Proceedings of NeurIPS 2024 Datasets and Benchmarks Track.*

Rebedea, T., Dinu, R., et al. (2023). NeMo Guardrails: A Toolkit for Controllable and Safe LLM Applications with Programmable Rails. *arXiv preprint arXiv:2310.10501.*

Inan, H., Upasani, K., et al. (2023). Llama Guard: LLM-based Input-Output Safeguard for Human-AI Conversations. *arXiv preprint arXiv:2312.06674*.

Bai, Y., Kadavath, S., et al. (2022). Constitutional AI: Harmlessness from AI Feedback. *arXiv preprint arXiv:2212.08073*.

Wei, A., Haghtalab, N., and Steinhardt, J. (2023). Jailbroken: How Does LLM Safety Training Fail? *Proceedings of NeurIPS 2023*.

Zou, A., Wang, Z., Kolter, Z., and Fredrikson, M. (2023). Universal and Transferable Adversarial Attacks on Aligned Language Models. *arXiv preprint arXiv:2307.15043*.

Khalil, H. K. (2002). *Nonlinear Systems*, Third Edition. Prentice Hall.

Wheeler, D. J. (2010). *Understanding Statistical Process Control*, Third Edition. SPC Press.